

Enron Submission Free-Response Questions

Wei Zhang

1. Summarize for us the goal of this project and how machine learning is useful in trying to accomplish it. As part of your answer, give some background on the dataset and how it can be used to answer the project question. Were there any outliers in the data when you got it, and how did you handle those? [relevant rubric items: “data exploration”, “outlier investigation”]

The goal for this project is to find the POI by machine learning.

The dataset is explored as shown in the picture:

```
The Enron data contains: 144
There are 18 POI
There are 50 NaN in feature salary
There are 106 NaN in feature deferral_payments
There are 21 NaN in feature total_payments
There are 141 NaN in feature loan_advances
There are 63 NaN in feature bonus
There are 127 NaN in feature restricted_stock_deferred
There are 96 NaN in feature deferred_income
There are 19 NaN in feature total_stock_value
There are 50 NaN in feature expenses
There are 43 NaN in feature exercised_stock_options
There are 53 NaN in feature other
There are 79 NaN in feature long_term_incentive
There are 35 NaN in feature restricted_stock
There are 128 NaN in feature director_fees
There are 58 NaN in feature to_messages
There are 58 NaN in feature from_poi_to_this_person
There are 58 NaN in feature from_messages
There are 58 NaN in feature from_this_person_to_poi
There are 58 NaN in feature shared_receipt_with_poi
There are 0 NaN in feature fraction_to_poi
There are 0 NaN in feature fraction_from_poi
```

The outliers are found in two ways:

1. By visualization: We can found the TOTAL outliers easily
2. Finding an outlier that is actually not a name, like THE TRAVEL AGENCY IN THE PARK

These two should be removed from the dataset.

There are extra data points looks like outliers, but they are POIs, so we should keep the data.

2. What features did you end up using in your POI identifier, and what selection process did you use to pick them? Did you have to do any scaling? Why or why not? As part of the assignment, you should attempt to engineer your own feature that does not come ready-made in the dataset -- explain what feature you tried to make, and the rationale behind it. (You do not necessarily have to use it in the final analysis, only engineer and test it.) In your feature selection step, if you used an algorithm like a decision tree, please also give the feature importances of the features that you use, and if you used an automated feature selection function like SelectKBest, please report the feature scores and reasons for your choice of parameter values. [relevant rubric items: “create new features”, “properly scale features”, “intelligently select feature”]

Before analyzing using model, I created two extra features, 'fraction_to_poi' and 'fraction_from_poi'. I have picked 21 features except 'email_address'.

During the exploration, I found out that there are negative values in some features like “exercised_stock_options”, “restricted_stock_deferred”. So MaxMinScaler is used to scale them into good format.

Different algorithms the SelectKBest has picked different combination of features, like during NB classification, I picked 6 features including:

```
The score of feature loan_advances is 30.7287746334
The score of feature poi is 15.858730906
The score of feature shared_receipt_with_poi is 15.8394424435
The score of feature from_this_person_to_poi is 10.7225708137
The score of feature deferred_income is 10.6338520484
The score of feature expenses is 9.68004143038
```

While in adaboost, I picked 12 features of them:

```
The score of feature loan_advances is 30.7287746334
The score of feature poi is 15.858730906
The score of feature shared_receipt_with_poi is 15.8394424435
The score of feature from_this_person_to_poi is 10.7225708137
The score of feature deferred_income is 10.6338520484
The score of feature expenses is 9.68004143038
The score of feature deferral_payments is 8.95913664769
The score of feature restricted_stock_deferred is 8.7922038527
The score of feature long_term_incentive is 8.05830631228
The score of feature other is 7.55511977732
The score of feature total_payments is 7.03793279819
The score of feature to_messages is 4.95866668397
```

3. What algorithm did you end up using? What other one(s) did you try? How did model performance differ between algorithms? [relevant rubric item: “pick an algorithm”]

I have tried three algorithm: GaussianNB(), AdaBoostClassifier(), KNeighbors_classifier()
The algorithm with best performance is GaussianNB().

```
This is a NB Classifier
The Training time: 0.002 s
The Predict time: 0.0 s
The accuracy: 0.907

=====
Pipeline(memory=None,
       steps=[('reduce_dim', SelectKBest(k=6, score_func=<function f_classif at 0x10df0d0c8>)), ('clf', GaussianNB(priors=None))])
  Accuracy: 0.84840      Precision: 0.41951      Recall: 0.35700 F1: 0.38
574  F2: 0.36797
    Total predictions: 15000      True positives: 714      False positives:
988  False negatives: 1286      True negatives: 12012
```

The stability of AdaBoostClassifier() is not stable, the performance is bellow:

```
This is a Adaboost Classifier
The Training time: 1.368 s
The Predict time: 0.0 s
The accuracy: 0.860

=====
Pipeline(memory=None,
       steps=[('reduce_dim', SelectKBest(k=13, score_func=<function f_classif at 0x10f3850c8>)), ('clf', AdaBoostClassifier(algorithm='SAMME',
       base_estimator=DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=None,
       max_features=None, max_leaf_nodes=None,
       ...one,
       splitter='best'),
       learning_rate=1.0, n_estimators=2, random_state=None))])
  Accuracy: 0.82407      Precision: 0.33385      Recall: 0.32100 F1: 0.32
730  F2: 0.32349
    Total predictions: 15000      True positives: 642      False positives:
1281 False negatives: 1358      True negatives: 11719
```

KNeighbors_classifier performed well in precision, but bad in recall:

```
This is a KNeighbors Classifier
The Training time: 8.158 s
The Predict time: 0.0 s
The accuracy: 0.860

=====
Pipeline(memory=None,
       steps=[('reduce_dim', SelectKBest(k=8, score_func=<function f_classif at 0x10a6ea0c8>)), ('clf', KNeighborsClassifier(algorithm='auto', leaf_size=5, metric='minkowski',
       metric_params=None, n_jobs=1, n_neighbors=2, p=2,
       weights='uniform'))])
  Accuracy: 0.86907      Precision: 0.53103      Recall: 0.15400 F1: 0.23
876  F2: 0.17949
    Total predictions: 15000      True positives: 308      False positives:
272  False negatives: 1692      True negatives: 12728
```

4. **What does it mean to tune the parameters of an algorithm, and what can happen if you don't do this well? How did you tune the parameters of your particular algorithm? (Some algorithms do not have parameters that you need to tune -- if this is the case for the one you picked, identify and briefly explain how you would have done it for the model that was not your final choice or a different model that does utilize parameter tuning, e.g. a decision tree classifier). [relevant rubric item: "tune the algorithm"]**

Take adaboost_classifier as example, there are three steps in modify the algorithm:

1. Selected good estimate as start, in this case, I chose the default estimate, decision_tree_classifier with min_samples_split = 5 in order to avoid over fitting.
2. Adjusted K of SelectKBest() to pick suitable features for analyzing. K usually set to be 11,12,13,14
3. Tune n_estimators to near 2 or 5 to specific the boundary.

5. **What is validation, and what's a classic mistake you can make if you do it wrong? How did you validate your analysis? [relevant rubric item: "validation strategy"]**

If I do it wrong, the true_positives value is always error, that means I didn't predict any poi is 1 properly, in other word, the model cannot recognize poi. So using tester.py, I will receive the message of Precision or recall may be undefined due to a lack of true positive predictions.

Then if the problem occurred, I should look back and see the accuracy of the algorithm is right. If the accuracy is right, then I should look at the features that I pick is properly chosen, or just tune the parameters of the algorithm.

6. **Give at least 2 evaluation metrics and your average performance for each of them. Explain an interpretation of your metrics that says something human-understandable about your algorithm's performance. [relevant rubric item: "usage of evaluation metrics"]**

Please refer the pictures I attached above.

GaussianNB() can estimate the poi with an accuracy of 0.85, positive predictive value of 0.41, and false omission rate of 0.36.

AdaBoostClassifier() have an average accuracy of 0.82, positive predictive value of 0.33, and false omission rate of 0.30.

So it is obvious that the GaussianNB() has better performance in both prediction and false prevention.